

CS221 Fall 2016 Project Final Report: Scrabble AI

Authors: Colleen Josephson {cajoseph} and Rebecca Greene {greenest}

Introduction

The goal of our project is to build an AI that plays Scrabble, a popular crossword board game in which players build words on a 15x15 board using tiles representing letters. Different letters have different values, and the number of points a player receives for forming a word is equal to the sum of the values of the tiles in the word, with multipliers depending on location on the board. There are a fixed set of tiles in the game, and at any point in the game, each player has access to at most 7 letter of these in their 'rack'. When tiles are placed on the board they are replaced with ones drawn at random from a bag. If the player uses all 7 tiles this is called a 'bingo' and the move receives a 50pt bonus, which can be about 1/8 of the total points in a tournament game.

Since each player can make no direct observations of the other player's rack, Scrabble can be considered a stochastic partially observable game [Russel and Norvig, 2003]. Thus it is very different from many games like Othello, chess, or Go in which each player has complete knowledge of the state of the game at any point. In addition, the incredible number of possible moves renders typical decision-tree models to be all but impossible.

In the following sections we will discuss the background and prior work, our approach to solving the problem, our experimental infrastructure, and an analysis of the results.

Background and Prior Work

A number of previous Scrabble AIs have been developed, starting with a Scrabble/crossword playing program developed by C. Sharpiron at the University of New York at Buffalo in 1982, and continuing today with Scrabble/crossword AIs like Quackle.

One of the most important publications in the history of Scrabble AIs was a paper by Andrew Appel and Guy Jacobson published in 1988 on 'The World's Fastest Scrabble Program'. In it, Appel and Jacobson describe the move generation algorithm that has since been used by all major Scrabble AIs. They convert the lexicon into a data structure called a 'directed acyclic word-graph' or 'dawg', that greatly expedites the search for valid moves. We use a simplified version of this algorithm, substituting dawgs for tries, which will be explained in greater detail below.

The most famous Scrabble AI was a program called 'Maven' that was developed by Brian Sheppard in 2002, which was the first Scrabble AI to beat human champion players. Maven divides the game Scrabble game into three subcategories- regular game play, the pre-end game, and the end-game. The 'end game' is the period after which all letters in the bag have been used. At this point the 'end game' is the period after which all letters in the bag have been used. At this point the game becomes completely observable, and a B* search on a game tree can be used to optimize agent game play. We based a lot of our implementation on the regular game play strategies used by Maven, but ignored the pre-end game and end-game

aspects of the game, as we were limited on time, and found the problem a less interesting pursuit. Maven is the AI currently used in official Hasbro Scrabble games.

The current champion of the Scrabble AIs is Quackle, developed by 5 researchers at MIT. Quackle uses a strategy that is very similar to Maven’s. Quackle is the current measuring stick against which all other Scrabble AIs are compared, and we analyze our AI’s performance against Quackle later in the paper.

Approach

Extensive vocabulary alone is not sufficient for a competitive Scrabble player. If a player optimizes for the best score on every turn they tend to retain tiles that are more difficult to use in play, leading to future racks that will produce a lower score. The best Scrabble players try to maintain their rack in a way that will be conducive to future high-scoring plays.

Computers can be preprogrammed with the entire permissible dictionary, since that dictionary is about 200,000 words long, and the 7 letters on the rack can be combined with those on the board in tens of thousands of different permutations, the search for possible moves becomes a nontrivial undertaking. Most Scrabble AIs give themselves a time limit to ensure reasonable progress of play.

Once a list of possible moves has been generated, the AI needs to select the best move as a weighted decision of the score that the move would generate, the opportunities it would provide of the other player, and the effect it would have on the future rack.

Scrabble AIs face the following challenges:

1. **Move generation:** create a list of possible moves from the state of the board, the letters in the rack, and the allowable words in the dictionary. This is a nontrivial search problem.
2. **Rack maintenance:** balance the tradeoff between getting the maximal score for a given turn with maintaining a rack that will be useful for future turns, which usually involves a weighted sum acquired with machine learning plus a number of Monte Carlo simulations.
3. **Adversarial gameplay:** avoid creating opportunities for the other player to place high scoring words

Once the bag has been emptied (all tiles are either on the board or in one of the racks), the game switches to being one with a completely known state. At this point evaluation techniques like minimax become useful to maximize score. This is commonly referred to as *endgame strategy*, and typically only state-of-the-art AIs go into this level of detail.

Model and Algorithms

Move Generation

To help solve the search problem, we’re using an algorithm created in the 1980s by Andrew Appel and Guy Jacobson. This algorithm remains the backbone of most competitive Scrabble

AI's today. Appel and Jacobson propose restructuring the Scrabble lexicon from a list of words into a trie or prefix tree, where each node is a partial word, the children of a node are words or partial words that can be created using that node (see Fig. 1). All terminal leaves of the trie are words, as are some interim nodes (e.g. 'dog' vs. 'dogs'), and the value of a node is a boolean indicating whether the string is a full word in the dictionary.

From Appel-Jacobsen '86




Figure 1: An Example of a Trie

To search for moves on the board, the algorithm examines all *anchors*, where an anchor is the space to the left (or above) an existing letter (for horizontal plays), or the space above an existing letter (see Fig. 2). Since a move in Scrabble must attach to an existing word (excepting the first move), this greatly reduces the 15x15 search space. After each new move on the board, the AI should update its list of anchor squares.




Figure 2: An Example of Anchor Squares

Before a potential move is added to the list of generated moves, it needs to be *crosschecked* to ensure that new strings formed in the orthogonal dimension also exist in the dictionary. For example, in Figure 3, when placing DOG underneath CATS, AD and TO are real words, but the vertical SG would fail a crosscheck. Since the crosscheck results for a given tile remain static unless the tiles adjacent to it change, the checks only need only be updated once per move, and only for the squares immediately adjacent to newly placed tiles.

The heart of the algorithm is a backtracking search with constraints that the final result must be a valid word (enforced by the structure of the trie itself), all crosschecks pass, and the new letters placed on the board come from the player's rack. The search algorithm has two recursive parts: ExtendLeft and ExtendRight. Below is pseudocode for the backtracking algorithm to place a horizontal word:

```

ExtendRight (PartialWord, node N, square):
  if square is not empty:
    if the letter l in square is an edge of N (PartialWord + l is a node):
      ExtendRight(PartialWord + l, N.children[l], nextSquare)
  else:
    if PartialWord is a word: LegalMoves.append(PartialWord)

```

```

for each letter l that is in rack and an edge out of N
    and in the cross-check set of square:
        remove l from the rack
        ExtendRight(PartialWord + l, N.children[l], nextSquare)
        put tile l back into the rack

```

ExtendLeft places tiles to the left of the anchor point and then calls ExtendRight. 'Limit' is the number of blank tiles between the current anchor point and the preceding anchorpoint (or the end of the board), capped at the rack size of 7.

```

ExtendLeft(PartialWord, node N, square, limit):
    ExtendRight(PartialWord, N, square)
    if limit > 0:
        for each letter l in rack that is an edge of N:
            remove l from the rack
            ExtendLeft(l+PartialWord, N' = N.children[l], nextSquare, limit -1)
            put tile l back into the rack

```

To generate a list of all legal moves for a given board, call LeftExtend("", root node, anchorSquare, Limit) on all anchors. See Figures 1 and 2 in 'Examples and Preliminary Data'

Rack Maintenance and Adversarial Gameplay

The next step of the problem is to try to decide which of the legal moves is best, given consideration of score, maintaining a reasonable rack, and not giving any advantages to the opponent. This is best done through a combination of Monte Carlo simulation and linear predictors with learned weights.

The raw score of a move is as a function of tile values and multipliers on the board, plus any resulting multi-word bonuses or bingos. This is simple to compute, but also a simplified view of the situation. A better metric is the agent/opponent point differential obtained as the result of a move, which is the result of the raw score of the move and also the opportunities it proves the opposing player. This differential is best computed through simulation. For each of the top raw-scoring moves, the AI runs a number of Monte Carlo simulations playing against itself with probable opponent racks. Since the turnover rate of racks is so high, and the computation required per move is rather extensive, most competitive AIs run their models with a search depth ≤ 3 . We use a depth-2 Monte Carlo simulation, essentially comparing the difference in score after both players make a move.

Many Scrabble AIs don't try to make any assumptions about the opposing player's rack, and just assign the opponent random unseen letters when running simulations. However, it is possible to use Bayes' algorithm to make a probabilistic model of the tiles a player had on their rack at the start of a turn given the move they made during that turn. For instance, if the opposing player used the letters 'C, T' to attach to an A and make "CAT", it is unlikely that they left an S on their rack, because otherwise they would have played "C,T,S" to make "CATS", which is a higher scoring word. More formally $P(\text{leave}|\text{play}) = \frac{P(\text{play}|\text{leave})P(\text{leave})}{P(\text{play})}$,

where $P(\text{leave})$ is the probability that certain tiles were left on the rack. When an AI's Monte Carlo simulations draw from this probability space rather than a random assignment, the algorithm has better performance on a level that is statistically significant (Richards and Amir, 2007). We didn't have time to successfully implement the Bayesian network to solve this final problem, and instead drew simulation racks at random from the list of unseen tiles, but we weighted the resulting score by the probability of that rack occurring.

Now that the scores have been computed by simulation, they are weighed along with heuristics for rack maintenance to select the best move to use. The features extracted for this purpose are:

1. whether or not a given tile is present in the rack
2. duplicated tiles, e.g. in the rack AABCDEF 'AA' would be a double
3. triples of a tile(CCC) for different letters
4. balance of vowels and constants (which has been proven to be an important factor in weight maintenance)
5. whether or not 'Q' has a corresponding 'U' in the rack at the same time

The weights for these features are learned by initially setting them all to zero, and having the AI play itself, running a stochastic gradient ascent since we want to *maximize* the score differential. Consequently the loss function we use is $\mathbf{w} \cdot \phi(A) - \mathbf{w} \cdot \phi(B)$. The SGA is implemented in `sga.py`.

Experimental Infrastructure

The two primary ways to test AIs are against humans or against other AIs. A common benchmark is whether or not an AI can beat a human world champion.

We did test our game against people, but collecting that data is time consuming and logistically challenging. To generate our data we played our AI against Quackle. We also played variants of our AI against itself. The variants we created are:

- **vanilla:** selects max scoring move from Appel-Jacobsen move generation algorithm
- **rack heuristic:** selects N moves with the top raw scores and creates a feature vector from the raw score and features extracted from the remaining rack, and applies a weight vector learned via 100-iteration stochastic gradient run. The move with the max weighted score is returned
- **Monte Carlo:** selects the top N scoring moves and simulates M different depth D future plays, and returns the score differentials for each move weighted by the probability of the opponent being able to make that move. For our results, we used $N=3$, $M=15$ and $D=2$.
- **Monte Carlo + rack heuristic:** includes the Monte Carlo score in the feature vector and applies a different weight vector; the move with the top weighted score is returned

Quackle vs CS221 Mode

Playing against Quackle in an automated manner proved to be a reasonably challenging task, as our code is in Python while Quackle was written in C++. To avoid re-implementing our AI in C++, we used a file-based system for interprocess communication. We modified the test system for Quackle in `quackle/test/testharness.cpp` to interface with our AI via reading and writing moves to files. We created one main method, `selfPlayCS221Game()`.

For the Python half, we wrote `'autorun.py'`, which starts the Quackle test harness and reads and writes to the shared files until the specified number of games has been played. The scores are stored in a dictionary, which is printed out at the end so the data can be processed in `Scrabble/images/plot.py`.

This was the most challenging part of our infrastructure to set up, but as a result we are able to see how our AI compares to world-class AI.

Self Play Mode

Playing our AI against itself is significantly more simple, and also provides good data on how much the different variants improve gameplay. Self-play mode is in `selfrun.py`, and does not require any file IO.

Human Mode

For testing various aspects of gameplay and to see how well the AI performed against average-ability humans, we created a human vs AI mode in `run.py`. Since human-generated data is time-consuming to collect, and neither the authors nor their friends describe themselves as skilled Scrabble players, we do not collect any statistics from human mode.

Results and Analysis

Overall, we are able to beat Quackle 11-13% of the time with an average score in the 280's. Since this is against a world class AI, we are quite pleased with the results. Figures 1-4 show the *score differential* between Quackle and our AI. The lower the differential, the better. For example, a differential of 20 means that Quackle won that game by 20 points.

Table 1 outlines additional statistics, such as the average scores of each AI, the average score differential, the number of games we lose by 2x points or more, and the maximum score ratio (a ratio of 1.5 means we scored 1.5 times as many points as Quackle for that game). Entries in Tables 1 and 2 marked with * were shorter runs due to time constraints. Table 1's MC+RH was 35 runs, and Table 2's MC and MC+RH were 80 and 53 runs respectively. Furthermore, the MC+RC weight vector was only run for 50 iterations instead of 100.

We were disappointed to see that the additional features like rack heuristics and Monte Carlo didn't make a big difference in game outcomes, at least while playing Quackle. We suspect that the gains given by rack heuristics and Monte Carlo are only noticeable when two AIs already perform comparably. The results of our selfplay games in Table 2 support this, as we win against our own AI a statistically significant amount more—a 62% win rate compared to roughly 50-50 for the control set.

	win rate	mean	opp mean	mean delta	2x loss rate	max win ratio
vanilla	11%	284.68	407.43	122.75	13%	1.358
RH	13%	288.07	410.0	121.93	11%	2.131
MC	13%	283.3	401.95	118.65	14%	1.42
MC w/ RH*	5%	285.38	412.22	126.83	13%	2.34

Table 1: Statistics of our AI variants against Quackle’s Speedy Player. *Mean delta* is the average difference between opponent and our score, *2x loss rate* is the rate at which the opponent scores 2x or higher, and *max win ratio* is the max ratio of our score to opponent’s.




Figure 1: Quackle vs CS221 Vanilla

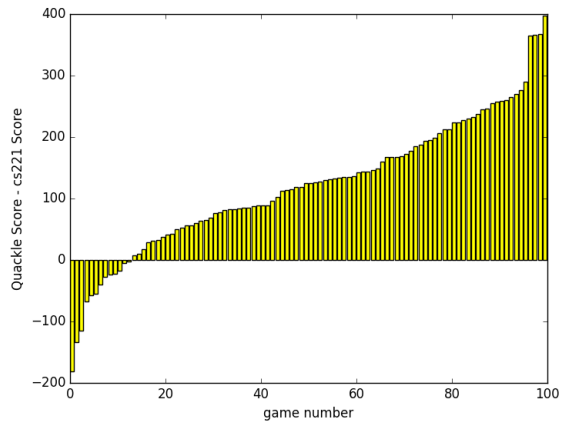


Figure 2: Quackle vs CS221 w/ Rack Heuristics

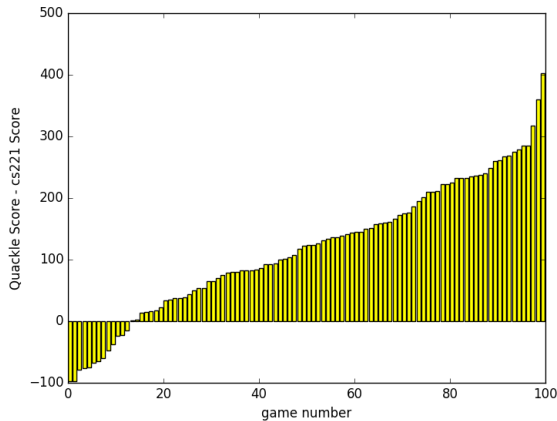


Figure 3: Quackle vs CS221 w/ Monte Carlo




Figure 4: Quackle vs CS221 with Monte Carlo and Rack Heuristics*

The rack heuristics perform similarly to the control. This is not surprising considering our weight vectors though (which can be found at the top of util.py). The weight for the raw score multiple orders of magnitude higher than any other feature. This means that the moves selected in the vanilla case and the rack heuristic case will not often differ. However, in the rack heuristic + Monte Carlo case the Monte Carlo score is only one order of magnitude smaller than the raw score, which means that the Monte Carlo score non-trivially impacts

which move is ultimately selected. This is evidenced by the increased win rate of 53%. Unfortunately the weight vector used on this dataset had a limited SGA run (50 iterations instead of 100). Also, since Monte Carlo games take 10x longer to play, we needed to terminate the data collection before it reached 100 games. We suspect that if we had more time the results would be even better.

	win rate	mean	opp mean	mean delta	2x loss rate	max win ratio
control	47%	354.24	359.25	-5.01	0%	1.811
RH	48%	351.59	353.81	-2.22	0%	1.509
MC*	62%	361.55	340.56	20.98	0%	1.41
MC w/ RH*	53%	364.62	347.07	17.54	0%	1.33

Table 2: Statistics of our AI playing against itself for 300 games. Control is the vanilla variant playing against itself, all other rows are variations playing against vanilla. For description of columns, see caption for Table 1.

One particular challenge to improving our performance was the long development cycle. For example, after adding or removing features to the feature extractor we need to re-run SGA and then re-run all datasets that depend on a weight vector. This is very time consuming, especially when we include Monte Carlo simulation, so we could only do a limited amount of algorithmic experimentation.

Finally, we suspect that a large part of the gap is simplifications we made to edge cases in gameplay. As mentioned earlier, we do not adapt our strategy for the end game. We also do not do detailed analysis for tile exchanges. We only exchange tiles if the Move Generator returns empty, and we do not strategize which tiles to exchange from our rack. Our wilcard behavior is also simplified—we simply choose a vowel if our rack currently has none, otherwise we choose a tile randomly. This is because a blank tile greatly increases the number of possible moves and non-trivially modifying the existing search algorithm.

Anecdotally, the AI does very well against average humans (the authors and their friends), even if it probably won’t beat a world champion.

Conclusion

Considering that we are both new to AI and had only 6 weeks to work on this project part-time, we are pretty pleased that we managed to beat the research-grade Quackle AI 13% of the time.

Neither of the authors will be playing Scrabble again anytime soon.

Code Appendix

Contents of Scrabble folder:

- agent.py: AI to choose an optimal move
- autorun.py: quackle vs CS221 AI mode
- brain.py: the Appel-Jacobsen move generator + Monte Carlo simulation
- baseline.py: baseline implementation from proposal
- board.py: the Scrabble board object
- boardTests.py: unit tests for board functionality
- brainTests.py: unit tests for move generator
- images/: images for our progress and final reports
- letterbag.py: the object that holds Scrabble tiles
- oracle.py: our oracle implementation
- pdfs/: pdfs of our proposal, progress, and poster
- pygtrie: the Google trie implementation we use for move generation
- results/: raw results, plots and statistics. Raw data stored and processed in plot.py
- run.py: human vs AI mode
- Scrabblewords.txt: our Scrabble dictionary
- selfrun.py: selfplay mode
- sga.py: implements stochastic gradient ascent
- treeBuilder.py: builds the trie from Scrabblewords.txt
- trie.p: trie pickle
- util.py: utility functions and variables

Human mode runs by simply executing `python run.py`. Selfplay runs via `python selfplay.py -n <NUM_GAMES>`. To use autorun.py, you need to build Quackle. Follow the README in the root of the quackle folder to build, and then go to quackle/test and run `qmake test.pro && make`

This will create an executable called 'test' in quackle/test. We'll be running the test program with these options: `./test --repetitions=N lexicon=cs221 --mode=cs221`

This tells the quackle test harness to run N games in cs221 mode with the cs221 lexicon. The code for cs221 mode is mainly defined in the function `selfPlayCS221Game()` inside `test/testharness.cpp`.

The Python half is in `Scrabble/autorun.py`. It interfaces with the quackle AI via text files in quackle/test. Quackle's data is written to `quackle/test/quacklegame-n.gcg`, which `autorun.py` reads and parses. After parsing and committing the quackle move, our AI calculates its own move, commits it, and writes it to the file `test/quackle/cs221game-n`.

To run Quackle vs. CS221:

```
python autorun.py -n <NUM_GAMES> -p <PATH_TO_QUACKLE_TEST>
```

Optionally, add a `-s` flag to suppress ASCII-board output, useful when running large batches.

References

- A.W. Appel, G.J. Jacobson, The worlds fastest Scrabble program, Comm. ACM 31 (5) (1988) 572578, 585.
- Mark Richards and Eyal Amir. Opponent Modeling in Scrabble. IJCAI-07., 14821487, 2007.
- Stuart Russell and Peter Norvig. Artificial Intelligence: A Modern Approach. Prentice-Hall, Englewood Cliffs, NJ, 2nd edition edition, 2003
- Brian Sheppard. World-championship caliber Scrabble. Artif. Intell., 134:241245, 2002.
- J. Katz-Brown, J. O’Laughlin, J. Fultz, M. Liberty, A. Buddhdev, Quackle,
<http://people.csail.mit.edu/jasonkb/quackle/>